

# Practical Machine Learning Course Project

//

*January 2016*

## The Weight Lifting Exercises Dataset Experiment

Six young health participants, were asked to perform one set of 10 repetitions of the Unilateral Dumbbell Biceps Curl in five different fashions: exactly according to the specification (Class A), throwing the elbows to the front (Class B), lifting the dumbbell only halfway (Class C), lowering the dumbbell only halfway (Class D) and throwing the hips to the front (Class E).

For data recording while performing the sets, the participants were wearing four 9 degrees of freedom Razor inertial measurement units (IMU), which provide three-axes acceleration, gyroscope and magnetometer. The sensors have been mounted in the users' glove, armband, lumbar belt and dumbbell.

The goal for the Weight Lifting Exercises experiment is to investigate "how (well)" an activity was performed by the wearer.

The goal of the project is to predict the manner in which they did the exercise. This is the "classe" variable in the data set.

Web site:

<http://groupware.les.inf.puc-rio.br/har> (<http://groupware.les.inf.puc-rio.br/har>).

Data Source, the Weight Lifting Exercises Dataset:

[http://groupware.les.inf.puc-rio.br/static/WLE/WearableComputing\\_weight\\_lifting\\_exercises\\_biceps\\_curl\\_variations.csv](http://groupware.les.inf.puc-rio.br/static/WLE/WearableComputing_weight_lifting_exercises_biceps_curl_variations.csv)  
([http://groupware.les.inf.puc-rio.br/static/WLE/WearableComputing\\_weight\\_lifting\\_exercises\\_biceps\\_curl\\_variations.csv](http://groupware.les.inf.puc-rio.br/static/WLE/WearableComputing_weight_lifting_exercises_biceps_curl_variations.csv))

Data Sources used for the Course Project:

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>  
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>)

The test data (for the quiz) are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>  
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>)

## Data Exploratory and Cleaning

```
#Reading the datasets, training and test
setwd("~/coursera/Practical Machine Learning")
#download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv",
"./pml-training.csv")
#download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv", "./
pml-testing.csv")
training<-read.csv("./pml-training.csv",stringsAsFactors = T,na.strings=c("NA", "", "#DI
V/0!"))
quiz<-read.csv("./pml-testing.csv",stringsAsFactors = T, na.strings=c("NA", "", "#DIV/
0!"))

#dimension of the Weight Lifting Exercises Dataset
dim(training)
```

```
## [1] 19622    160
```

```
#Number of observations for each participant and exercise
table(training$user_name,training$classe)
```

```
##
##           A      B      C      D      E
## adelmo    1165   776   750   515   686
## carlitos   834   690   493   486   609
## charles    899   745   539   642   711
## eurico     865   592   489   582   542
## jeremy    1177   489   652   522   562
## pedro      640   505   499   469   497
```

The feature statistics (average, mn, max , var, etc...) are calculated for each window and have missing values for most of the observations, so they are removed from the analysis.

Note that for a specific participant and time of execution (timestamp variable) you know which exercise has been performed, also knowing the window number automatically tell you which type of exercise has been performed and by whom. Since the goal is to identify the type of exercise performed based on sensors measurements then the identification variables are also removed from the data set.

```
#Data cleaning
#Keep only non empty features in the testing file
allvar <- names(training)
vBeg <- c("kurtosis_", "skewness_", "max_", "min_", "amplitude_", "var_", "avg_", "stddev_")
va <- vector()
for(i in 1:length(vBeg)){ va <- c(va,grep(vBeg[i],allvar,value = FALSE))}
training <- training[,-c(1:7,va)]
quiz <- quiz[,-c(1:7,va)]
```

```
#Load all the required libraries and suppression of warning messages
suppressWarnings(suppressMessages(library(caret, quietly=T)))
suppressWarnings(suppressMessages(library(rpart, quietly=T)))
suppressWarnings(suppressMessages(library(MASS, quietly=T)))
suppressWarnings(suppressMessages(library(kernlab, quietly=T)))
suppressWarnings(suppressMessages(library(gbm, quietly=T)))
suppressWarnings(suppressMessages(library(plyr, quietly=T)))
suppressWarnings(suppressMessages(library(randomForest, quietly=T)))
```

Split of the Weight Lifting Exercises Dataset in 2 parts, a simple splitting based on the outcome (classe). The first part (train) is used to explore the data and fit the models, the second (test) is used to evaluate the out of sample accuracy of each model.

```
# create training set indexes with 70% of data
set.seed(12345)
inTrain <- createDataPartition(y=training$classe,p=0.7, list=FALSE)
train <- training[inTrain,]
test <- training[-inTrain,]
dim(train)
```

```
## [1] 13737    53
```

```
dim(test)
```

```
## [1] 5885    53
```

Identification of highly correlated predictors, some models may benefit from reducing the level of correlation between the predictors. Only one predictor variable will be kept from the group of correlated variables. It will also significantly reduce the time execution to fit the models.

```
highlyCorDescr <- findCorrelation(cor(train[, -length(train)]), cutoff = .8)
train<-train[-highlyCorDescr[2:length(highlyCorDescr)]]
test<-test[-highlyCorDescr[2:length(highlyCorDescr)]]
quiz<-quiz[-highlyCorDescr[2:length(highlyCorDescr)]]
```

Use of the nearZeroVar function to identify the variables that have no variability, these variables are not useful when we want to construct a prediction model.

```
# print nearZeroVar table
near0<-nearZeroVar(train,saveMetrics=TRUE)
near0
```

```
##          freqRatio percentUnique zeroVar  nzv
## yaw_belt      1.022161    12.92130742  FALSE FALSE
## total_accel_belt 1.054462     0.19654946  FALSE FALSE
## gyros_belt_x    1.033024     0.96090850  FALSE FALSE
## gyros_belt_y    1.129367     0.48773386  FALSE FALSE
## gyros_belt_z    1.085784     1.18657640  FALSE FALSE
## accel_belt_z    1.103110     2.13292568  FALSE FALSE
## magnet_belt_x   1.069231     2.18388294  FALSE FALSE
## magnet_belt_y   1.123894     2.08196841  FALSE FALSE
## magnet_belt_z   1.024845     3.18118949  FALSE FALSE
## roll_arm       55.488372    17.76224794  FALSE FALSE
## pitch_arm      91.807692    20.13540074  FALSE FALSE
## yaw_arm        33.138889    19.29096600  FALSE FALSE
## total_accel_arm 1.032206     0.47317464  FALSE FALSE
## gyros_arm_y     1.470588     2.66433719  FALSE FALSE
## gyros_arm_z     1.203390     1.70342870  FALSE FALSE
## accel_arm_y     1.163265     3.77083788  FALSE FALSE
## accel_arm_z     1.030612     5.58346073  FALSE FALSE
## magnet_arm_x    1.016129     9.55812768  FALSE FALSE
## magnet_arm_z    1.061728     9.13591032  FALSE FALSE
## roll_dumbbell   1.155556    87.00589648  FALSE FALSE
## pitch_dumbbell  2.048077    84.92392808  FALSE FALSE
## yaw_dumbbell    1.155556    86.36529082  FALSE FALSE
## total_accel_dumbbell 1.080559     0.31302322  FALSE FALSE
## gyros_dumbbell_y 1.207229     1.94365582  FALSE FALSE
## accel_dumbbell_y 1.023121     3.28310403  FALSE FALSE
## magnet_dumbbell_x 1.181818     7.84741938  FALSE FALSE
## magnet_dumbbell_y 1.333333     5.98383927  FALSE FALSE
## magnet_dumbbell_z 1.084034     4.78998326  FALSE FALSE
## roll_forearm    12.013100    13.59831113  FALSE FALSE
## pitch_forearm   62.477273    18.83235059  FALSE FALSE
## yaw_forearm     15.531073    12.88490937  FALSE FALSE
## total_accel_forearm 1.146991     0.50229308  FALSE FALSE
## gyros_forearm_x 1.016173     2.02373153  FALSE FALSE
## gyros_forearm_z 1.154762     2.15476450  FALSE FALSE
## accel_forearm_x 1.111111     5.67081604  FALSE FALSE
## accel_forearm_y 1.029412     7.11945840  FALSE FALSE
## accel_forearm_z 1.009174     4.08386111  FALSE FALSE
## magnet_forearm_x 1.075472    10.50447696  FALSE FALSE
## magnet_forearm_y 1.000000    13.27800830  FALSE FALSE
## magnet_forearm_z 1.095238    11.71289219  FALSE FALSE
## classe         1.469526     0.03639805  FALSE FALSE
```

```
dim(train)
```

```
## [1] 13737    41
```

No predictor has only one distinct value or has a near zero variance, so the 41 variables from that list are those who will be used to build our prediction model.

## Methodology

For the course project, 5 machine learning classification methods will be compared:

- \* Classification Tree, CART (rpart)
- \* Linear Discriminant Analysis (lda)
- \* Least Squares Support Vector Machine with Radial Basis Function Kernel (lssvmRadial)
- \* Stochastic Gradient Boosting (gbm)
- \* Random Forest (rf)

The models are fitted on the training part of the data set (first split) using a cross validation procedure with k=10 folds, the default cross validation method of the caret train function. The out of sample error rate will be measured on the independent test data set (second split) for all the model.

The model with the best accuracy on the test data set will be the chosen one to answer the quiz test.

- Classification Tree, CART (rpart)

```
#rpart model
set.seed(12345)
modFitrpart <- train(classe ~ ., method = "rpart", data = train, trControl = trainControl(
  method = "cv"))
```

```
#rpart model fit
print(modFitrpart)
```

```
## CART
##
## 13737 samples
##    40 predictor
##      5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 12361, 12364, 12363, 12364, 12364, 12365, ...
## Resampling results across tuning parameters:
##
##  cp          Accuracy   Kappa      Accuracy SD   Kappa SD
##  0.01963178  0.5778536  0.4691779  0.02236006   0.02690560
##  0.02776930  0.5343899  0.4106954  0.02351234   0.03898975
##  0.03707659  0.3807065  0.1660515  0.10461144   0.17807823
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.01963178.
```

```

pred_rpart<-predict(modFitrpart,newdata=test[-length(test)])
#rpart out of sample accuracy
cmrpart<-confusionMatrix(pred_rpart, test$classe)
cmrpart$overall['Accuracy']

```

```

## Accuracy
## 0.5682243

```

- Linear Discriminant Analysis (lda)

```

#modFitlda
set.seed(12345)
modFitlda <- train(classe ~ ., method = "lda", data = train, trControl = trainControl(method = "cv"))

```

```

#modFitlda model fit
print(modFitlda)

```

```

## Linear Discriminant Analysis
##
## 13737 samples
##    40 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 12361, 12364, 12363, 12364, 12364, 12365, ...
## Resampling results
##
## Accuracy   Kappa      Accuracy SD   Kappa SD
## 0.6575665   0.5670318   0.01068907    0.01348488
##
##

```

```

pred_lda<-predict(modFitlda,newdata=test[-length(test)])
#lda out of sample accuracy
cmlda<-confusionMatrix(pred_lda, test$classe)
cmlda$overall['Accuracy']

```

```

## Accuracy
## 0.6484282

```

- Least Squares Support Vector Machine with Radial Basis Function Kernel (lssvmRadial)

```
#modFitsvm
set.seed(12345)
modFitsvm <- train(classe ~ ., method = "lssvmRadial", data = train, trControl = trainControl(method = "cv"))
```

```
#modFitsvm model fit
print(modFitsvm)
```

```
## Least Squares Support Vector Machine with Radial Basis Function Kernel
##
## 13737 samples
## 40 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 12361, 12364, 12363, 12364, 12364, 12365, ...
## Resampling results
##
## Accuracy Kappa Accuracy SD Kappa SD
## 0.7991016 0.7451986 0.04042991 0.05144062
##
## Tuning parameter 'sigma' was held constant at a value of 0.01757531
##
```

```
pred_svm<-predict(modFitsvm,newdata=test[-length(test)])
#lssvmRadial out of sample accuracy
cmsvm<-confusionMatrix(pred_svm, test$classe)
cmsvm$overall['Accuracy']
```

```
## Accuracy
## 0.8346644
```

- Stochastic Gradient Boosting (gbm)

```
#modFitgbm
set.seed(12345)
modFitgbm <- train(classe ~ ., method = "gbm", data = train, trControl = trainControl(method = "cv"), verbose=FALSE)
```

```
#modFitgbm model fit
print(modFitgbm)
```

```
## Stochastic Gradient Boosting
##
## 13737 samples
##    40 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 12361, 12364, 12363, 12364, 12364, 12365, ...
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  Accuracy   Kappa     Accuracy SD
##  1                   50      0.7248309  0.6507813  0.014716447
##  1                   100     0.7907137  0.7350301  0.009763385
##  1                   150     0.8313352  0.7865311  0.008739113
##  2                   50      0.8335167  0.7889911  0.010583483
##  2                   100     0.8909545  0.8619683  0.008412107
##  2                   150     0.9188346  0.8972863  0.007381244
##  3                   50      0.8788702  0.8466366  0.012351644
##  3                   100     0.9296079  0.9109173  0.008169796
##  3                   150     0.9502099  0.9369875  0.007566751
##  Kappa SD
##  0.018959554
##  0.012499172
##  0.011247408
##  0.013463758
##  0.010717224
##  0.009373944
##  0.015730235
##  0.010389390
##  0.009596822
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 150,
##  interaction.depth = 3, shrinkage = 0.1 and n.minobsinnode = 10.
```

```
pred_gbm<-predict(modFitgbm,newdata=test[-length(test)])
#gbm out of sample accuracy
cmgbm<-confusionMatrix(pred_gbm, test$classe)
cmgbm$overall['Accuracy']
```

```
## Accuracy
## 0.9500425
```

- Random Forest (rf)



```
#modFitrF
set.seed(12345)
modFitrF <- train(classe ~ ., method = "rf", data = train, trControl = trainControl(metho
d = "cv"))
```

```
#modFitrF model fit
print(modFitrF)
```

```
## Random Forest
##
## 13737 samples
## 40 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 12361, 12364, 12363, 12364, 12364, 12365, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.9912632 0.9889469 0.002147138 0.002717251
## 21 0.9898814 0.9871988 0.002236130 0.002829110
## 40 0.9852949 0.9813951 0.002741541 0.003470233
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
pred_rf<-predict(modFitrF,newdata=test[-length(test)])
#rf out of sample accuracy
cmrf<-confusionMatrix(pred_rf, test$classe)
cmrf$overall['Accuracy']
```

```
## Accuracy
## 0.9870858
```

```
#rf expected out of sample error
1-cmrf$overall['Accuracy']
```

```
## Accuracy
## 0.01291419
```

The machine learning method with the best accuracy on the test data set is the random forest (0.9871). This is the one choose for the quiz prediction. The expected out of sample error of the random forest model is 1.3%.

```
#Quiz prediction with Random Forest Model  
predict(modFitrfr,newdata=quiz[-length(test)])
```

```
## [1] B A B A A E D B A A B C B A E E A B B B  
## Levels: A B C D E
```

The score obtained for the quizz is 20/20.